



## ECE 3057: Architecture, Concurrency and Energy in Computation Fall 2019

# Lecture 2: Instruction Set Architecture

David Anderson

School of Electrical and Computer Engineering  
Georgia Institute of Technology

anderson@gatech.edu

Acknowledgment: Slides adapted from MIT 6.823 (Arvind and J. Emer).  
CMU 18-447 (O. Mutlu) and GT ECE 3057 (S. Yalamanchili, T. Krishna)

## 1823: Charles Babbage

A polynomial can be computed from difference tables

An example:

$$\begin{aligned} f(n) &= n^2 + n + 41 \\ d_1(n) &= f(n) - f(n-1) = 2n \\ d_2(n) &= d_1(n) - d_1(n-1) = 2 \end{aligned}$$

$$\begin{aligned} d_2(n) &= 2 \\ d_1(n) &= d_1(n-1) + d_2(n) \\ f(n) &= f(n-1) + d_1(n) \end{aligned}$$



Charles Babbage 1791-1871  
Lucasian Professor of Mathematics,  
Cambridge University, 1827-1839  
**Father of the "Computer"**

n	0	1	2	3	4 ...
$d_2(n)$		2	2	2	2
$d_1(n)$		2	4	6	8
$f(n)$	41	43	47	53	61

What computation unit needed?  
*an adder!*

3

## Topics

- Why do we need Instruction Set Architectures?
  - A Historical Perspective
- Elements of an ISA
- MIPS ISA

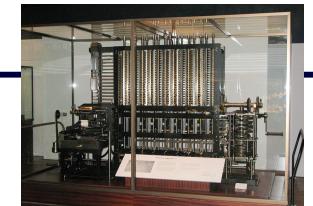
ECE 3057 | Summer 2019 | L02: ISA

David Anderson, School of ECE, Georgia Tech

May 16, 2019

4

## Difference Engine



1823

- Babbage's paper is published

1834

- The paper is read by Scheutz & his son in Sweden

1842

- Babbage gives up the idea of building it;  
He is onto Analytic Engine!

1855

- Scheutz displays his machine at the Paris World Fair
- Can compute any 6th degree polynomial
- Speed: 33 to 44 32-digit numbers per minute!
  - Scheutz machine is at the Smithsonian.
  - Difference Engine #2 – London science museum
  - copy at computer museum in CA

ECE 3057 | Summer 2019 | L02: ISA

David Anderson, School of ECE, Georgia Tech

May 16, 2019

## Analytic Engine

The forerunner to the modern digital computer

5

- 1833: Babbage's paper was published
  - conceived during a hiatus in the development of the difference engine
- Inspiration: Jacquard Looms
  - looms were controlled by punched cards
    - The set of cards with fixed punched holes dictated the pattern of weave => **program**
    - The same set of cards could be used with different colored threads => **inputs**
- 1871: Babbage dies
  - The machine remains unrealized
    - It is not clear if the analytic engine could be built even today using only mechanical technology

ECE 3057 | Summer 2019 | L02: ISA

David Anderson, School of ECE, Georgia Tech

May 16, 2019



Ada Lovelace, 1815-1852  
**The first programmer**

## Ada Byron aka Lady Lovelace

- Translated lectures of Luigi Menabrea who published notes of Babbage's lectures in Italy
- Lovelace considerably embellished notes and described how **Analytical Engine** could be used to calculate Bernoulli numbers.
- The first program!

## From Mechanical to Electrical

7

- **Harvard Mark I (1944)**
  - Proposed by Howard Aiken – Professor of Physics at Harvard and built by IBM Endicott laboratories
  - Essentially mechanical but had some electro-magnetically controlled relays and gears
    - Weighed 5 tons and had 750,000 components
    - A synchronizing clock that beat every 0.015 seconds
  - Performance
    - 0.3 sec for addition, 6 seconds for Multiplication, 1 minute for sine
    - But -- broke down once a week!
- **ENIAC**
  - Designed and built by Eckert and Mauchly at the University of Pennsylvania during 1943-45
  - **The first, completely electronic, operational, general-purpose analytical calculator!**
    - 30 tons, 72 square meters, 200KW
    - Lights dimmed in Philadelphia when turned on.
  - Application: Ballistic Calculations
  - Performance
    - Read in 120 cards per minute, Addition 200  $\mu$ s, Division 6 ms
    - 1000 times faster than Mark I
    - But frequent tube burnouts reduced accuracy

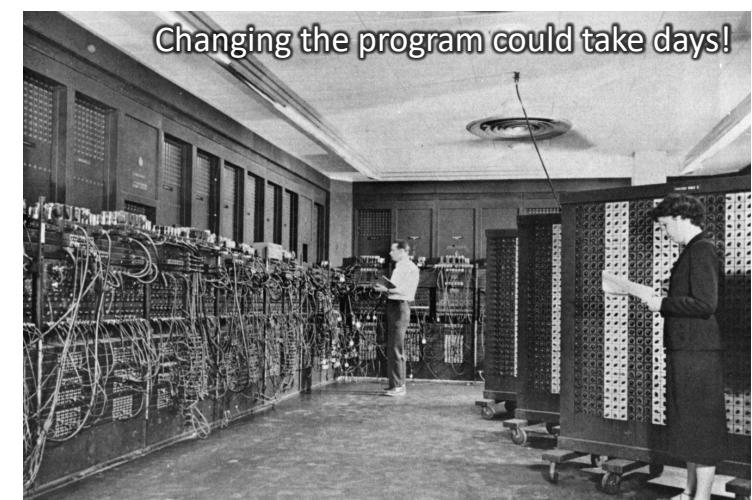
ECE 3057 | Summer 2019 | L02: ISA

David Anderson, School of ECE, Georgia Tech

May 16, 2019

## Electronic Numerical Integrator and Computer (ENIAC)

8



ECE 3057 | Summer 2019 | L02: ISA

David Anderson, School of ECE, Georgia Tech

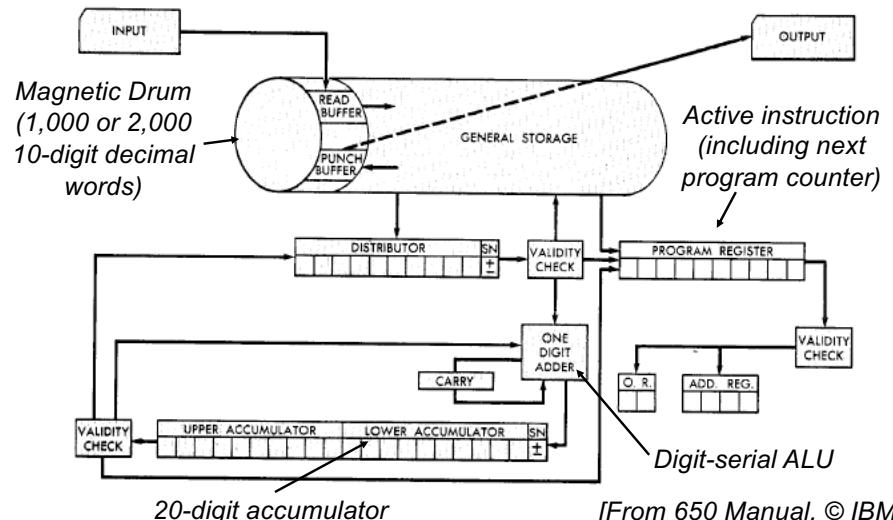
May 16, 2019

# Electronic Discrete Variable Automatic Computer (EDVAC)

9

- ENIAC's programming system was external
  - Sequences of instructions were executed independently of the results of the calculation
  - Human intervention required to take instructions "out of order"
- EDVAC was designed by Eckert, Mauchly and von Neumann in 1944 to solve this problem
  - Solution was the **stored program computer**
    - "program can be manipulated as data"
    - The same storage can be used to store program and data
- First Draft of a report on EDVAC was published in 1945, but just had von Neumann's signature!
  - Without a doubt the most influential paper in computer architecture**

# The IBM 650 (1953-54)



[From 650 Manual, © IBM]

## Programmer's view of IBM 650

11

A drum machine with 44 instructions

Instruction: 60 1234 1009

"Load the contents of location 1234 into the distribution; put it also into the upper accumulator; set lower accumulator to zero; and then go to location 1009 for the next instruction."

- Programmer's view of the machine was inseparable from the actual hardware implementation
- Good programmers optimized the placement of instructions on the drum to reduce latency!

## Compatibility Problem at IBM

12

By early 60's, IBM had 4 incompatible lines of computers!

701	→	7094
650	→	7074
702	→	7080
1401	→	7010

Each system had its own

- Instruction set
- I/O system and Secondary Storage:  
magnetic tapes, drums and disks
- assemblers, compilers, libraries,...
- market niche  
business, scientific, real time, ...

⇒ IBM 360

# IBM 360 : Design Premises

Amdahl, Blaauw and Brooks, 1964

13

## Architecture of the IBM System/360

**Abstract:** The architecture\* of the newly announced IBM System/360 features four innovations:

1. An approach to storage which permits and exploits very large capacities, hierarchies of speeds, read-only storage for microprogram control, flexible storage protection, and simple program relocation.
2. An input/output system offering new degrees of concurrent operation, compatible channel operation, data rates approaching 5,000,000 characters/second, integrated design of hardware and software, a new low-cost, multiple-channel package sharing main-frame hardware, new provisions for device status information, and a standard channel interface between central processing unit and input/output devices.
3. A truly general-purpose machine organization offering new supervisory facilities, powerful logical processing operations, and a wide variety of data formats.
4. Strict upward and downward machine-language compatibility over a line of six models having a performance range factor of 50.

\*The term **architecture** here is used to describe the attributes of a system as seen by a programmer, i.e., the conceptual structure and functional behavior, as distinct from the organization of the dataflow and controls, the logical design and the physical **implementation**

ECE 3057 | Summer 2019 | L02: ISA

David Anderson, School of ECE, Georgia Tech

May 16, 2019

## IBM 360: Initial Implementations (1964)

	Model 30	...	Model 70
Memory Capacity	8K - 64 KB		256K - 512 KB
Memory Cycle	2.0µs	...	1.0µs
Datapath	8-bit		64-bit
Circuit Delay	30 nsec/level		5 nsec/level
Registers	in Main Store		in Transistor
Control Store	Read only 1µsec		Dedicated circuits

- Six implementations (Models, 30, 40, 50, 60, 62, 70)
- 50X performance difference cross models
- ISA completely hid the underlying technological differences between various models.
- With minor modifications, IBM 360 ISA is still in use!

ECE 3057 | Summer 2019 | L02: ISA

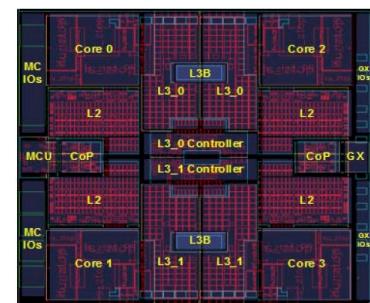
David Anderson, School of ECE, Georgia Tech

May 16, 2019

## IBM 360: Forty-Six years later... zEnterprise196 Microprocessor

15

- 1.4 billion transistors, Quad core design
  - Up to 96 cores (80 visible to OS) in one multichip module
  - 5.2 GHz, IBM 45nm SOI CMOS technology
    - 64-bit virtual addressing
  - original 360 was 24-bit; 370 was a 31-bit extension
  - Superscalar, out-of-order
    - Up to 72 instructions in flight
  - Variable length instruction pipeline: 15-17 stages
  - Each core has 2 integer units, 2 load-store units and 2 floating point units
  - 8K-entry Branch Target Buffer
    - Very large buffer to support commercial workload
  - Four Levels of caches:
    - 64KB L1 I-cache, 128KB L1 D-cache
    - 1.5MB L2 cache per core
    - 24MB shared on-chip L3 cache
    - 192MB shared off-chip L4 cache



[ September 2010 ]

ECE 3057 | Summer 2019 | L02: ISA

David Anderson, School of ECE, Georgia Tech

May 16, 2019

## Elements of an ISA

16

ECE 3057 | Summer 2019 | L02: ISA

David Anderson, School of ECE, Georgia Tech

May 16, 2019

## Elements of an ISA

- Instructions
  - Opcodes, Addressing Modes, Data Types
  - Instruction Types and Formats
  - Registers, Condition Codes
- Memory
  - Address space, Addressability, Alignment
  - Virtual memory management
- Call, Interrupt/Exception Handling
- Access Control, Priority/Privilege
- I/O: memory-mapped vs. instr.
- Task/thread Management
- Power and Thermal Management
- Multi-threading support, Multiprocessor support



Intel® 64 and IA-32 Architectures  
Software Developer's Manual  
Volume 1:  
Basic Architecture

## Two Kinds of ISAs

### ■ Complex instruction Set Computing (CISC)

- An instruction does a lot of work, e.g. many operations
  - Insert in a doubly linked list
  - Compute FFT
  - String copy

### ■ Reduced Instruction Set Computing (RISC)

- An instruction does small amount of work, it is a primitive
  - Add
  - XOR
  - Shift

## Complex vs Simple Instructions

- Advantages of Complex instructions
  - Denser encoding → smaller code size → better memory utilization, saves off-chip bandwidth, better cache hit rate (better packing of instructions)
  - Easier for assembly-language programming
  - Simpler compiler: no need to optimize small instructions as much (with caveats)
- Disadvantages of Complex Instructions
  - Larger chunks of work → compiler has less opportunity to optimize
  - More complex hardware → translation from a high level to control signals and optimization needs to be done by hardware
  - Compiler is limited in fine-grained optimizations it can do

## Examples

- CISC
  - X86, PDP-11, VAX, ...
- RISC
  - Alpha, MIPS, ARM, RISC-V

### ■ Most Popular ISAs Today?

- x86 vs ARM
  - x86 micro-ops → RISC-like micro instructions
  - RISC-V - an open-source ISA from Berkeley

## Microarchitecture

- Implementation of an ISA under specific **design constraints and goals**
- **Constraints**
  - Cost
  - Performance
  - Maximum power consumption
  - Energy consumption (battery life)
  - Availability
  - Reliability and Correctness
  - Time to Market
- **Goals**
  - determined by the “Problem” space (application space)
  - the intended users/market
    - E.g., smartphone vs laptop vs HPC vs cloud

## Microarchitecture

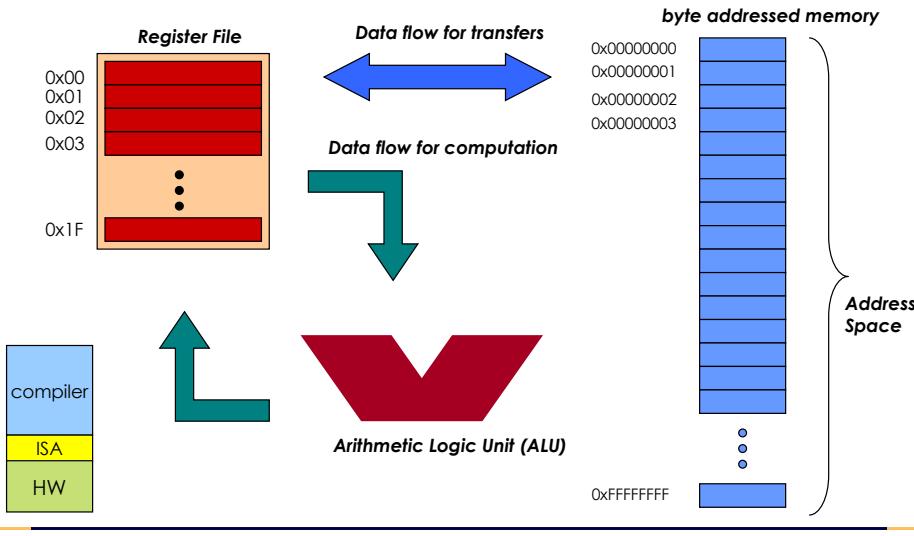
- Anything done in hardware without exposure to software
- Pipelining
- In-order versus out-of-order instruction execution
- Memory access scheduling policy
- Speculative execution
- Superscalar processing (multiple instruction issue?)
- Clock gating
- Caching? Levels, size, associativity, replacement policy
- Prefetching?
- Voltage/frequency scaling?
- Error correction?

## MIPS ISA

## Readings

- Useful Reference listing all MIPS instructions and encodings:
  - <http://www.mrc.uidaho.edu/mrc/people/jff/digital/MIPSir.html>
- Chapter 2
  - 2.1, Figure 2.1, 2.2 – 2.7
  - 2.9, Figure 2.15 (2.9.1 in online text)
  - 2.10, 2.16, 2.17
- Appendix A9, A10

## Instruction Set Architecture (ISA)



ECE 3057 | Summer 2019 | L02: ISA

David Anderson, School of ECE, Georgia Tech

May 16, 2019

## MIPS ISA

- Arithmetic and Logical Instructions
- Review memory organization
- Memory (data movement) instructions
- Control flow instructions
- Procedure/Function calls
- Program assembly, linking, & encoding

Not being covered

## MIPS ISA

R-type	<table border="1"><tr><td>op</td><td>rs</td><td>rt</td><td>rd</td><td>shamt</td><td>funct</td></tr></table>	op	rs	rt	rd	shamt	funct
op	rs	rt	rd	shamt	funct		
Arithmetic instruction format							
I-type	<table border="1"><tr><td>op</td><td>rs</td><td>rt</td><td colspan="3">address/immediate</td></tr></table>	op	rs	rt	address/immediate		
op	rs	rt	address/immediate				
Transfer, branch, immediate.							
J-type	<table border="1"><tr><td>op</td><td colspan="5">target address</td></tr></table>	op	target address				
op	target address						
Jump instruction							
Field size	<table border="1"><tr><td>6 bits</td><td>5bits</td><td>5bits</td><td>5bits</td><td>5bits</td><td>6 bits</td></tr></table>	6 bits	5bits	5bits	5bits	5bits	6 bits
6 bits	5bits	5bits	5bits	5bits	6 bits		

Useful Reference listing all MIPS instructions and encodings:  
<http://www.mrc.uidaho.edu/mrc/people/jff/digital/MIPSir.html>

ECE 3057 | Summer 2019 | L02: ISA

David Anderson, School of ECE, Georgia Tech

May 16, 2019

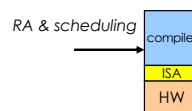
## MIPS ISA

- Arithmetic and Logical Instructions
- Review memory organization
- Memory (data movement) instructions
- Control flow instructions
- Procedure/Function calls
- Program assembly, linking, & encoding

May 16, 2019

## Arithmetic and Logical Instructions

- Operands must be registers
  - only 32 registers provided
- All memory accesses accomplished via loads and stores to/from registers
  - A common feature of RISC processors
- Compiler associates variables with registers
  - Some registers are not visible to the programmer
    - Program counter
    - Status register
    - .....



## MIPS Programmer Visible Registers

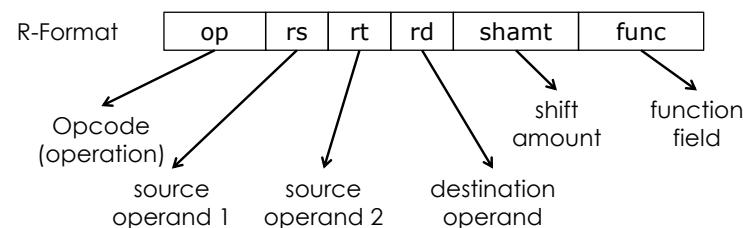
Register	Names	Usage by Software Convention
\$0	\$zero	Hardwired to zero
\$1	\$at	Reserved by assembler
\$2 - \$3	\$v0 - \$v1	Function return result registers
\$4 - \$7	\$a0 - \$a3	Function passing argument value registers
\$8 - \$15	\$t0 - \$t7	Temporary registers, caller saved
\$16 - \$23	\$s0 - \$s7	Saved registers, callee saved
\$24 - \$25	\$t8 - \$t9	Temporary registers, caller saved
\$26 - \$27	\$k0 - \$k1	Reserved for OS kernel
\$28	\$gp	Global pointer
\$29	\$sp	Stack pointer
\$30	\$fp	Frame pointer
\$31	\$ra	Return address (pushed by call instruction)
\$hi	\$hi	High result register (remainder/div, high word/mult)
\$lo	\$lo	Low result register (quotient/div, low word/mult)

## Logical Operations

- Instructions for bitwise manipulation
- Useful for extracting and inserting groups of bits in a word
- e.g.: masking, inverting, ...

Operation	C	Java	MIPS
Shift left	<<	<<	sll
Shift right	>>	>>>	srl
Bitwise AND	&	&	and, andi
Bitwise OR			or, ori
Bitwise NOT	~	~	nor

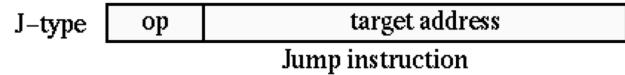
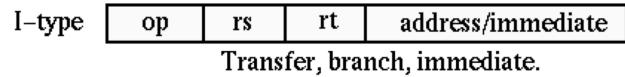
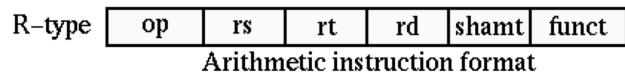
## Encoding: Instruction Format



- Instructions, like registers and words of data, are also 32 bits long
  - Example: add \$t0, \$s1, \$s2
  - registers have numbers, \$t0=9, \$s1=17, \$s2=18

Opcodes on page A-50 (Figure 7.10.2)  
Encodings – Section A10 (7.10)

## MIPS ISA

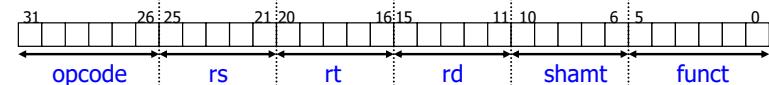


Field size

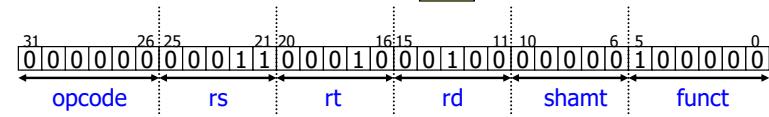
6 bits	5bits	5bits	5bits	5bits	6 bits
--------	-------	-------	-------	-------	--------

Useful Reference listing all MIPS instructions and encodings:  
<http://www.mrc.uidaho.edu/mrc/people/jff/digital/MIPSir.html>

## R-Type Encoding Example: ADD

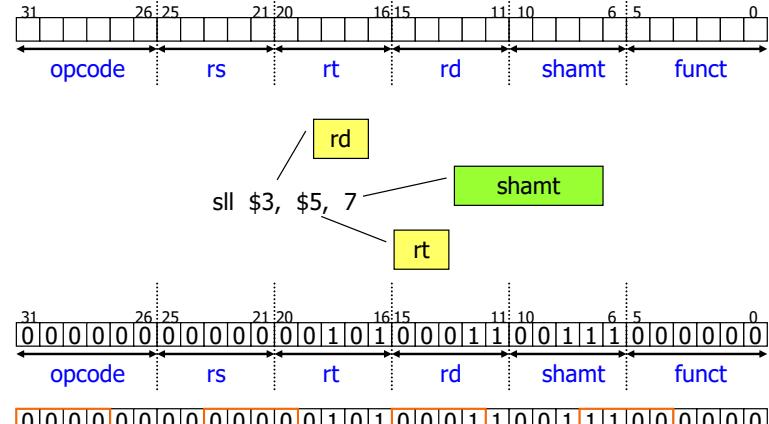


add \$4, \$3, \$2



Encoding = 0x00622020

## R-Type Encoding Example: SLL



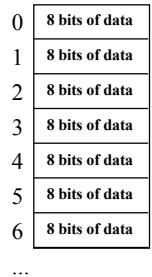
Encoding = 0x000519C0

## MIPS ISA

- Arithmetic and Logical Instructions
- Review memory organization
- Memory (data movement) instructions
- Control flow instructions
- Procedure/Function calls
- Program assembly, linking, & encoding

## Memory Organization

- Viewed as a large, single-dimension array, with an address.
- A memory address is an index into the array
- "Byte addressing" means that the index points to a byte of memory.



## Memory Organization

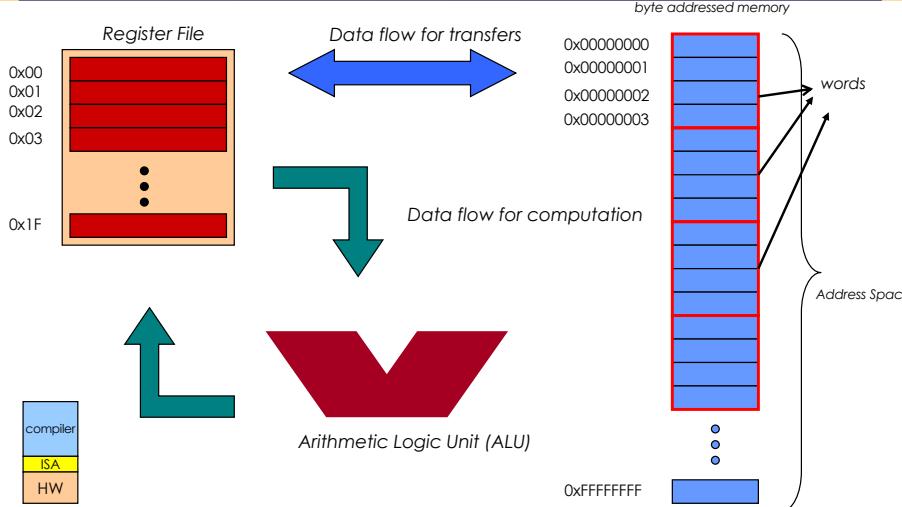
- Bytes are nice, but most data items use larger "words"
  - MIPS provides LW/LB and SW/SB instructions
- For MIPS, a word is 32 bits or 4 bytes.

0	32 bits of data
4	32 bits of data
8	32 bits of data
12	32 bits of data

Registers hold 32 bits of data

- $2^{32}$  bytes with byte addresses from 0 to  $2^{32}-1$
- $2^{30}$  words with byte addresses 0, 4, 8, ...  $2^{32}-4$
- Words are aligned
  - i.e., what are the least 2 significant bits of a word address?

## Instruction Set Architecture (ISA)

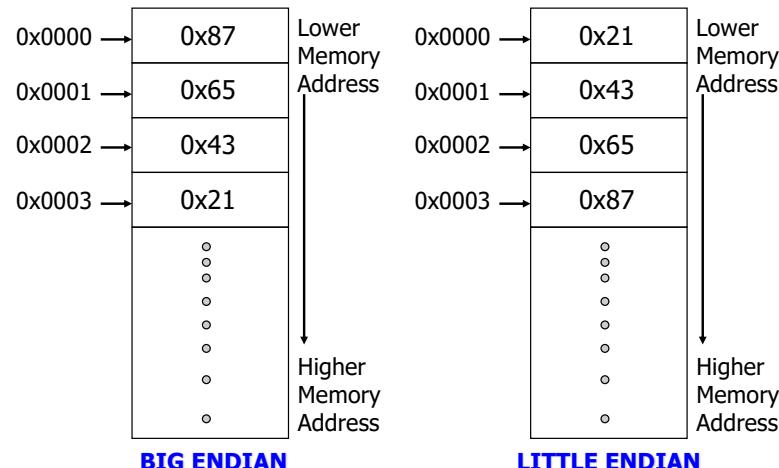


## Endianness [defined by Danny Cohen 1981]

- Byte ordering — How is a multiple byte data word stored in memory
- Endianness (from Gulliver's Travels)
  - Big Endian**
    - Most significant byte of a multi-byte word is stored at the lowest memory address
    - e.g. Sun Sparc, PowerPC
  - Little Endian**
    - Least significant byte of a multi-byte word is stored at the lowest memory address
    - e.g. Intel x86
- Some embedded & DSP processors would support both for interoperability
  - Most MIPS architectures (except R8000) supported both

## Example of Endian

Store 0x87654321 at address 0x0000, byte-addressable



## MIPS ISA

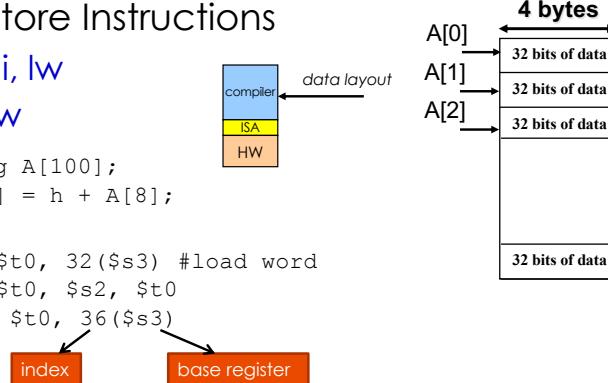
- Arithmetic and Logical Instructions
- Review memory organization
- **Memory (data movement) instructions**
- Control flow instructions
- Procedure/Function calls
- Program assembly, linking, & encoding

## Memory Instructions

- Load and Store Instructions
- Load: lb, lui, lw
- Store: sb, sw

C code: long A[100];  
A[9] = h + A[8];

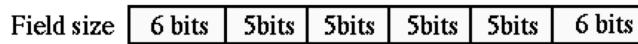
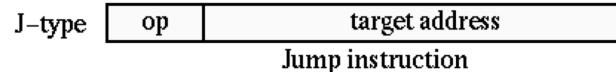
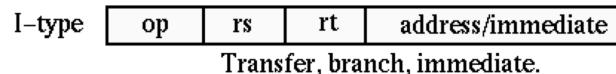
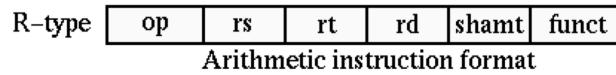
MIPS code: lw \$t0, 32(\$s3) #load word  
add \$t0, \$s2, \$t0  
sw \$t0, 36(\$s3)



## Registers vs Memory

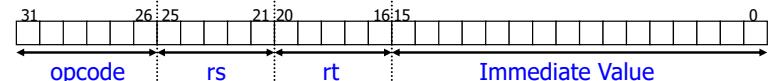
- Registers are faster to access than memory
  - Why? Smaller is faster
  - Rationale for the Memory Hierarchy
- Operating on memory data requires loads and stores
  - More instructions to be executed
- Compiler must use registers for variables as much as possible
  - Only spill to memory for less frequently used variables
  - Register usage optimization is important!

## MIPS ISA

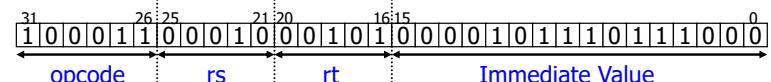


Useful Reference listing all MIPS instructions and encodings:  
<http://www.mrc.uidaho.edu/mrc/people/jff/digital/MIPSir.html>

## I-Type Encoding Example: LW



lw \$5, 3000(\$2)

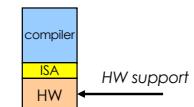


10001100010001010010000101110000

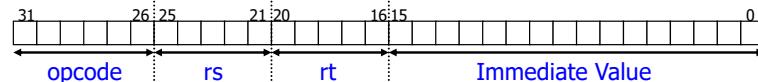
Encoding = 0x8C450BB8

## What are Immediate values?

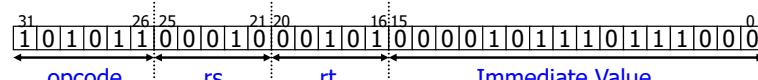
- Small constants are used quite frequently (50% of operands)
  - e.g., A = A + 5;  
B = B + 1;  
C = C - 18;
- Solutions?
  - put 'typical constants' in memory and load them.
  - create hard-wired registers (like \$zero) for constants like one.
  - Use immediate values
- Make Common Case Fast!
  - Small constants (<16-bit) are common
  - Immediate operand avoids a load instruction



## I-Type Encoding Example: SW



sw \$5, 3000(\$2)



10101100010001010010000101110000

Encoding = 0xAC450BB8

## Immediate Operands

- MIPS Instructions:

```
addi $29, $29, 4
slti $8, $18, 10
andi $29, $29, 6
ori $29, $29, 4
```

- No subtract immediate instruction

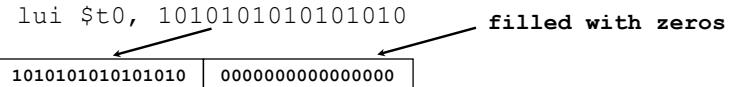
- Just use a negative constant  
`addi $s2, $s1, -1`

- Hardwired values useful for common operations

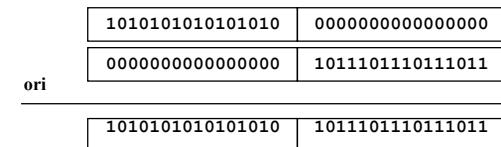
- E.g., move between registers  
`add $t2, $s1, $zero`

## What about larger constants?

- We'd like to be able to load a 32 bit constant into a register
- Must use two instructions, new "load upper immediate" instruction



- Then must get the lower order bits right, i.e.,  
`ori $t0, $t0, 1011101110111011`



## Pseudo Instructions

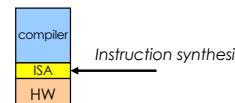
- Assembly instructions implemented with a mix of MIPS instructions

`li $t0, 0xAAAA BBBB`

=> `lui $t0, 0xAAAA`  
`ori $t0, $t0, 0xBBBB`

`move $t1, $t2`

=> `add $t1, $t2, $zero`



## MIPS ISA

- Arithmetic and Logical Instructions
- Review memory organization
- Memory (data movement) instructions
- Control flow instructions**
- Procedure/Function calls
- Program assembly, linking, & encoding

## Control

- Decision making instructions
  - alter the control flow,
  - i.e., change the "next" instruction to be executed
- MIPS conditional branch instructions:
 

```
bne $t0, $t1, Label
beq $t0, $t1, Label
```
- Example: if (i==j) h = i + j;
 

```
bne $s0, $s1, Label
add $s3, $s0, $s1
Label: ....
```

## Control

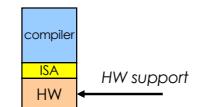
- MIPS unconditional branch instructions:
 

```
j label
```
- Example:
 

```
if (i!=j)
    h=i+j;
else
    h=i-j;
Lab1: sub $s3, $s4, $s5
Lab2:...
```

Assembler calculates address

- Can you build a simple loop?



## Compiling Loop Statements

- C code:
 

```
while (save[i] == k) i += 1;
    ■ i in $s3, k in $s5, address of save in $s6
```

- Compiled MIPS code:
 

```
Loop: sll $t1, $s3, 2 #multiply by 4
        add $t1, $t1, $s6
        lw $t0, 0($t1)
        bne $t0, $s5, Exit
        addi $s3, $s3, 1
        j Loop
Exit: ...
```

## Encoding: Branches & Jumps

- Instructions:
 

```
bne $t4,$t5,Label
beq $t4,$t5,Label
j Label
```

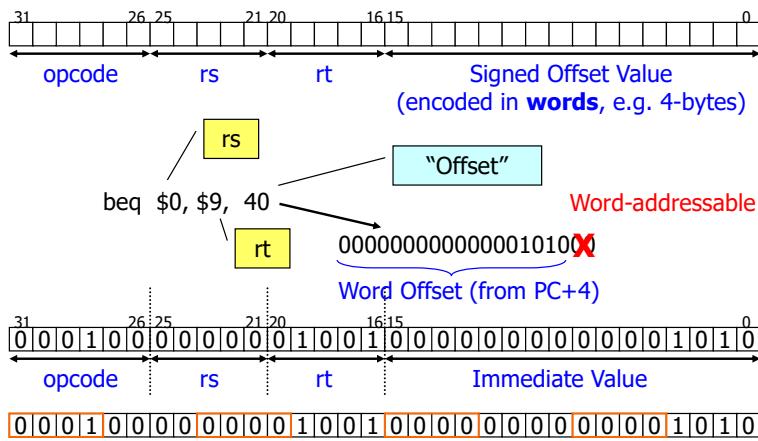
Next instruction is at Label if \$t4 ≠ \$t5  
Next instruction is at Label if \$t4 = \$t5  
Next instruction is at Label

- Formats:
 

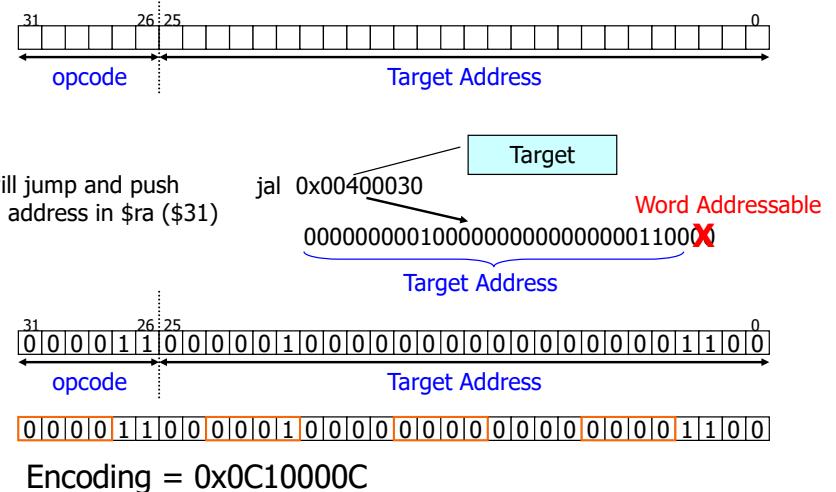
*Opcodes on page A-50 (Figure 7.10.2)  
Encodings – Section A10 (7.10)*

I	op	rs	rt	16 bit address
J	op			26 bit address

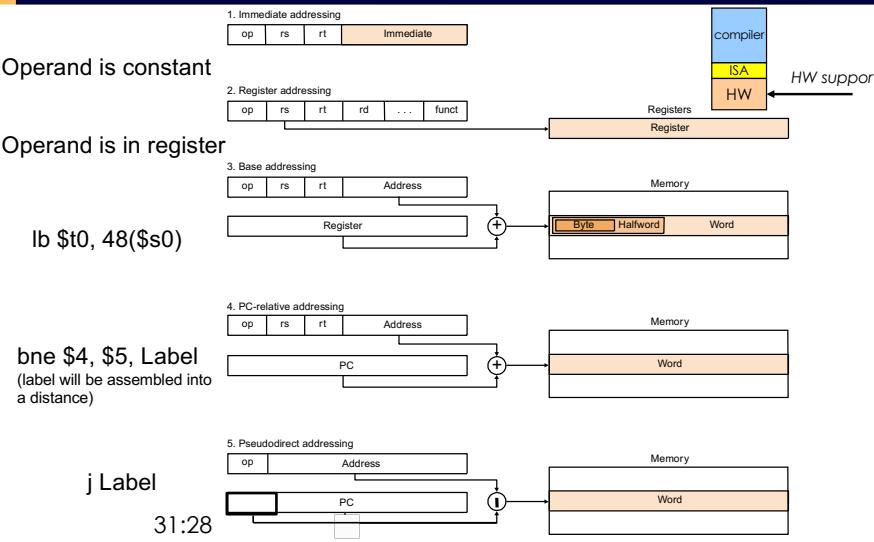
## I-Type Encoding Example: BEQ/BNE



## J-Type Encoding Example: JAL



## Addressing Modes



## How far can you branch?

- BEQ/BNE
  - $2^{16}$  words relative to PC+4
    - Most branches are local
  - What if branch is too far to encode?
    - Assembler rewrites code
- JR/JAL
  - $2^{26}$  words within boundary starting from PC[31:28]

## Summary: MIPS ISA

- Simple instructions all 32 bits wide

- Very structured

- Only three instruction formats

R	op	rs	rt	rd	shamt	funct
I	op	rs	rt	16 bit address		
J	op			26 bit address		

- Rely on compiler to achieve performance

*Opcodes on page A-50 (Figure 7.10.2)*

*Encodings – Section A10 (7.10)*