



ECE 3057: Architecture, Concurrency and Energy in Computation Summer 2019

http://tusharkrishna.ece.gatech.edu/teaching/uca_f18/

Lecture 08: Speculation

David Anderson

School of Electrical and Computer Engineering
Georgia Institute of Technology

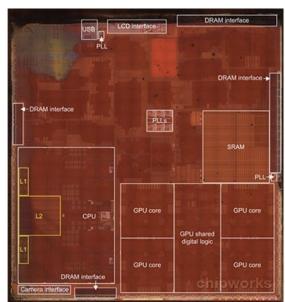
Acknowledgment: Lecture slides adapted from GT ECE 3056 (T. Krishna and S. Yalamanchilli), MIT 6.823 (Arvind and J. Emer)

Speculation in Practice

Apple found in infringement of University of Wisconsin CPU patent, faces \$862M in damages

By Mikey Campbell
Tuesday, October 13, 2015, 02:24 pm PT (05:24 pm ET)

A Wisconsin jury on Tuesday found Apple in infringement of computer processor patent owned by the University of Wisconsin Alumni Research Foundation, a ruling that could potentially lead to a damages payout of \$862.4 million.



Inside look at Apple's A7 | Source: Chipworks

U.S. Patent No. 5,781,752 for a "Table based data speculation circuit for parallel processing computer"

This "table" is used by Apple A7 and A8 → iPhone 5S, iPad Air, iPad Mini, iPhone 6, iPhone 6 Plus ...
aka **address speculation**

3

Resolving Data/Control Hazards

- Strategy 0: **remove dependence**
 - Detect and eliminate the dependence at the software level
- Strategy 1: **stall**
 - Wait for the result to be available by freezing earlier pipeline stages
- Strategy 2: **bypass/forward**
 - Route data as soon as possible after it is calculated to the earlier pipeline stage
- Strategy 3: **speculate**
 - Guess the result and proceed
 - Guessed correctly → no special action required
 - Guessed incorrectly → kill and restart
- Strategy 4: **do something else**
 - Switch to some other task/thread while dependence resolves

ECE 3057 | Su 2019 | L08: Speculation

David Anderson, School of ECE, Georgia Tech

June 4, 2019

2

When can we speculate?

- Data Hazards
 - An instruction depends on the value produced by a previous instruction
 - **Can we speculate?**
 - Maybe – if we can guess the output value of an ALU or LW instruction
- Control Hazards
 - Next instruction's PC depends on the previous instruction
 - For which instructions is this true?
 - Branches and Jumps
 - **Can we speculate?**
 - Yes – this lecture

ECE 3057 | Su 2019 | L08: Speculation

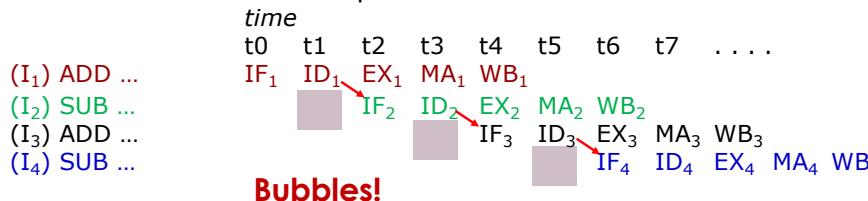
David Anderson, School of ECE, Georgia Tech

June 4, 2019

4

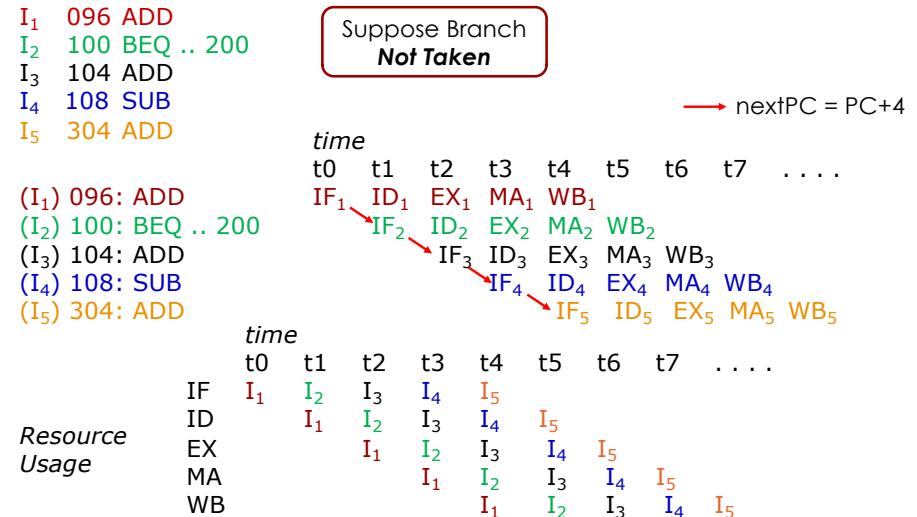
PC Calculation Bubbles

- In which stage do we know the opcode of the current instruction?
 - Decode
- Suppose we wait to do PC calculation till we know what the opcode is



- What is a good guess for next PC?
 - PC+4

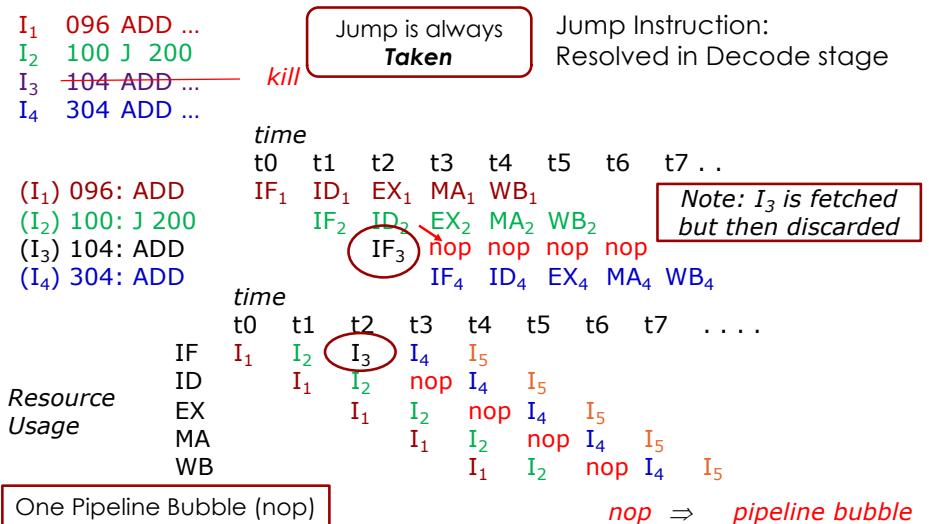
Example of correct speculation



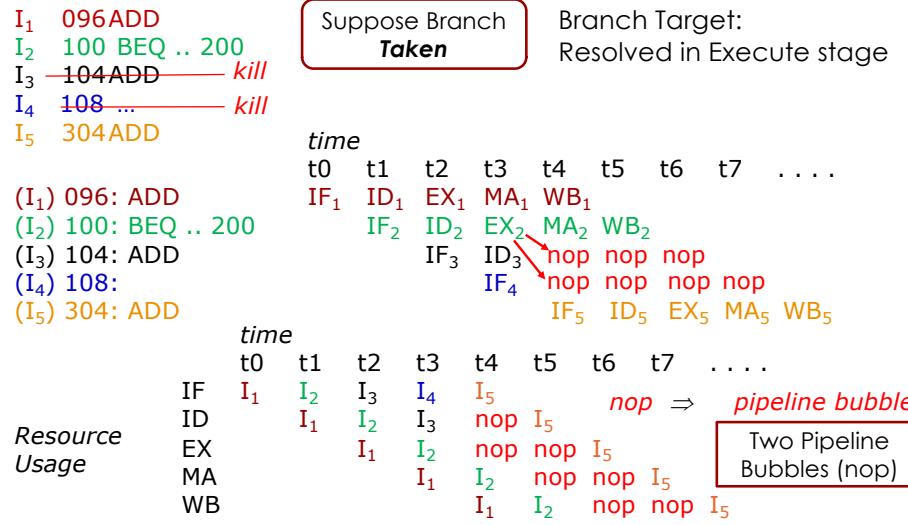
What if we mis-speculate?

- When do we know we have mis-specified?
 - Jump Instruction
 - At Decode stage
 - Branch Instruction
 - At Execute stage
- What do we do upon mis-speculation?
 - Start fetching instructions from the correct PC
 - What happens to the instructions we have already fetched?
 - Will be "flushed"
 - How?
 - Insert NOP in place of speculated instruction
 - Shows up as bubble in the pipeline

Example of Mis-speculation: Jump



Example of Mis-speculation: Branch



Why are branches important?

- Branches are very frequent
 - Approx. 20% of all instructions, i.e., 1 in every 5
 - Suppose 75% of them are Taken
 - $CPI = 1 + 0.2 \cdot 75 \cdot 2 = 1.3 \rightarrow 30\% \text{ slowdown}$
 - Worse penalty with deeper pipelines

- If we stall till we resolve the branch
 - Branch outcome is known after B cycles
 - B bubbles in the pipeline
 - B very large for long pipelines

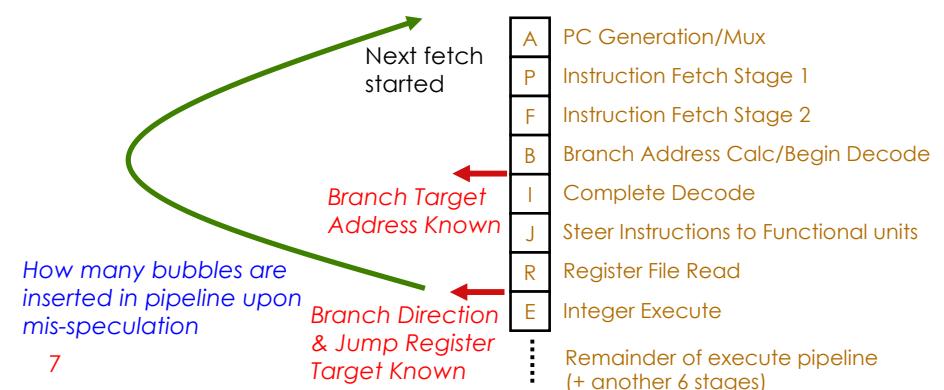
Mis-prediction penalty

- 5-stage Pipeline
 - Jumps: 1-cycle
 - Branches: 2-cycle

- What about deeper pipelines?

Branch Mis-prediction Penalty in Modern Pipelines

UltraSPARC-III instruction fetch pipeline stages
(in-order issue, 4-way superscalar, 750MHz, 2000)

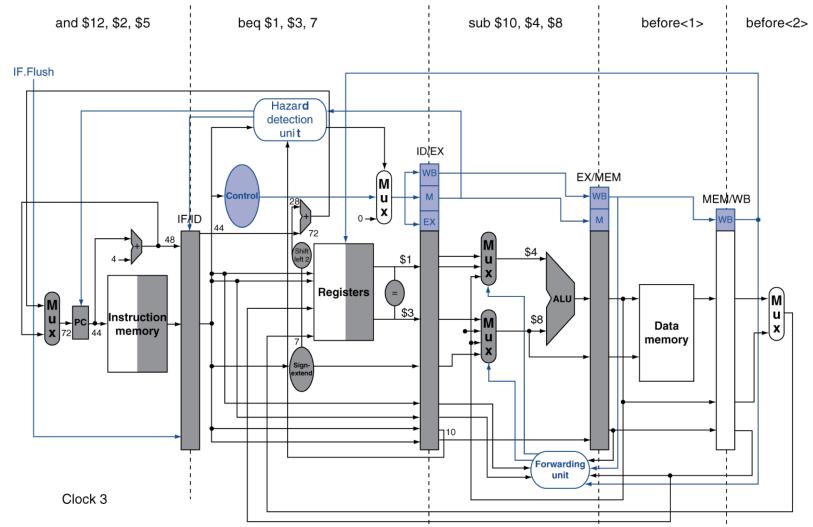


Reducing Branch Mis-prediction Penalty in 5-stage Pipeline

13

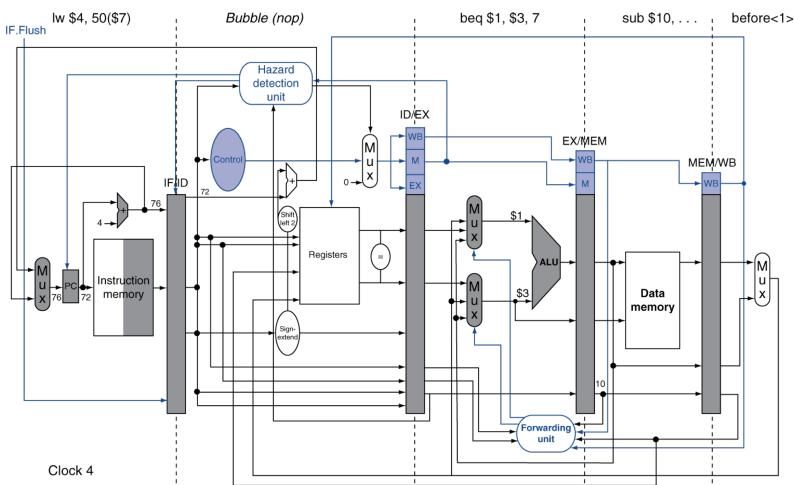
- Register Comparator
 - Add hardware to determine branch outcome in ID stage
 - Compare values of the two registers read
 - Now branches are identical to jumps → 1-cycle penalty
- Delay Slots (in MIPS)
 - Guarantee that the instruction after a Branch is always executed without being flushed
 - How?
 - Compiler adds a useful instruction or a NOP in the delay slot
 - Challenge?
 - Too specific to the microarchitecture

Example: Branch Taken



Chapter 4 — The Processor — 14

Example: Branch Taken



Chapter 4 — The Processor — 15

Improving Branch Prediction Accuracy

- What does $\text{nextPC} = \text{PC} + 4$ mean for branch instructions?
 - Predicting that the branch is Not Taken
- How true is that for real applications?
 - Consider the following piece of code


```
for (i=0; i < 100000; i++)
{
    C[i] = A[i] + B[i]
}
```

↓

// Suppose \$R0 = 0, and \$R10 is the iteration value

```
0xDC08    LW
...
0xDC48    SUBI R10, R10, 1
0xDC4C    BNE R10, R0, -0x44
```

Branch Instruction

What will be the most common outcome of this branch? **Taken**

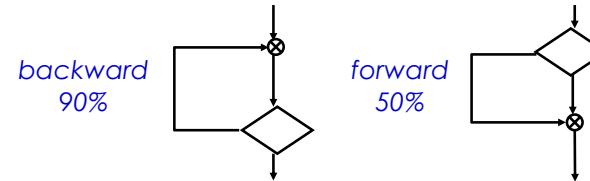
16

Branch Prediction

- Each instruction fetch depends on one or two pieces of information from the preceding instruction:
 - Is the preceding instruction a **taken branch**?
 - If so, what is the **target address**?

Static Branch Prediction

Overall probability a branch is taken is ~60-70% but:



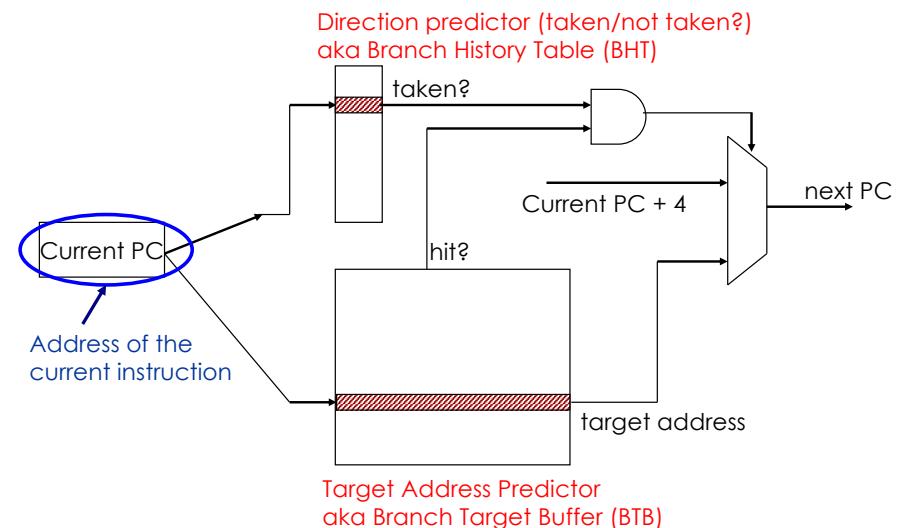
ISA can attach preferred direction semantics to branches, e.g., Motorola MC88110
bne0 (preferred taken) beq0 (not taken)

ISA can allow arbitrary choice of statically predicted direction,
e.g., HP PA-RISC, Intel IA-64
typically reported as ~80% accurate

Dynamic Branch Prediction

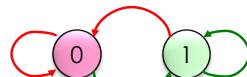
- Temporal correlation**
 - The way a branch resolves may be a good predictor of the way it will resolve at the next execution
 - Example: loops
- Spatial correlation**
 - Several branches may resolve in a highly correlated manner
 - Example: a preferred path of execution

Branch Predictor



One-Bit Direction Predictor

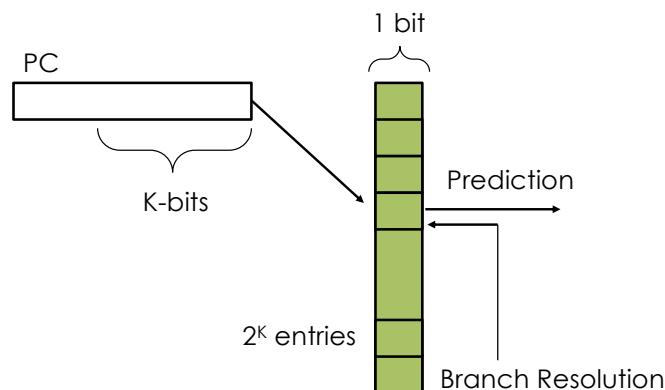
For each branch, remember its last outcome (T/NT)



Finite State Machine for
1-bit counter
(aka Last-time Predictor)

- Predict NT
- Predict T
- Transition on T outcome
- Transition on NT outcome

One-Bit Predictor Implementation



How well does this perform on loop branches?

One-bit predictor performance

Suppose the following loop is called multiple times

0xDC08:	<code>for (i=0; i < 100000; i++)</code>	Misprediction Rate
	{	2/100000 = 0.002 %

}

i	0	1	2	...	99998	99999	0	1	...	99998	99999	0	1	...	99998	99999	0	1
Predict	N	T	T	...	T	T	N	T	...	T	T	N	T	...	T	T	N	T
Actual	T	T	T	...	T	N	T	T	...	T	N	T	T	...	T	N	T	T

How well does this perform on loop branches?

Mispredicts **twice** for every use of loop.

One-bit predictor performance

0xDC08:	<code>for (i=0; i < 100000; i++)</code>	Misprediction Rate
	{	2/100000 = 0.002 %

0xDC44:	<code>if ((i % 100) == 0)</code>	<code>CountHundred();</code>	2000/100000 = 2 %
----------------	----------------------------------	-------------------------------	-------------------

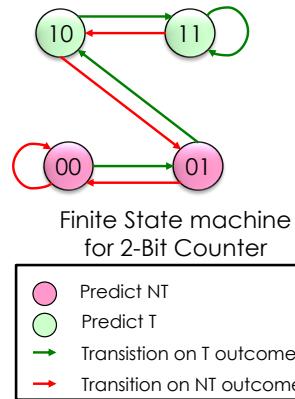
0xDC50:	<code>if ((i & 1) == 1)</code>	<code>CountOddNums();</code>	100000/100000 = 100 %
----------------	------------------------------------	-------------------------------	-----------------------

}

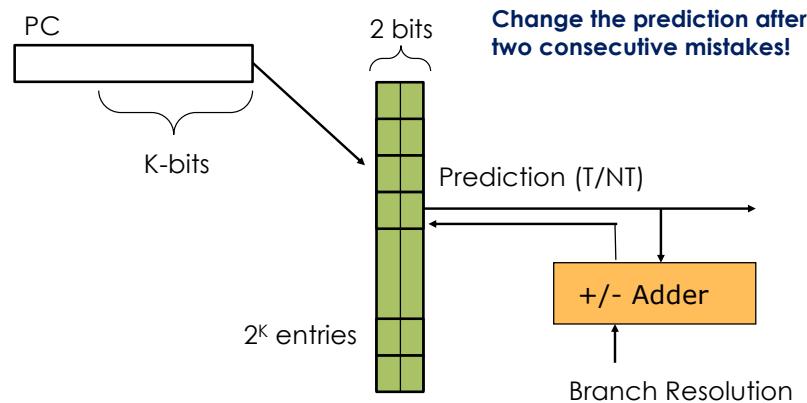
What happens on loop branches?

Mispredicts **twice** for every use of loop.

Two-bit Direction Predictor



Two-bit Predictor Implementation



Two-bit Predictor Performance

0xDC08: `for (i=0; i < 100000; i++)` Misprediction Rate
 { $1/100000 = 0.001\%$

}

i	0	1	2	...	99998	99999	0	1	...	99998	99999	0	1	...	99998	99999	0	1
State	00	01	10	11	11	11	10	11	11	11	11	10	11	11	11	11	10	11
Predict	N	N	T	...	T	T	T	T	...	T	T	T	T	...	T	T	T	T
Actual	T	T	T	...	T	N	T	T	...	T	N	T	T	...	T	N	T	T

What happens on loop branches?

Mispredicts **one** for every use of loop.

Two-bit Predictor Performance

0xDC08: `for (i=0; i < 100000; i++)` Misprediction Rate
 { $1/100000 = 0.001\%$

0xDC44: `if ((i % 100) == 0)` CountHundred(); $1000/100000 = 1\%$

0xDC50: `if ((i & 1) == 1)` CountOddNums(); ?

0xDC50 Predictor Starting State:

00 → 00 → 01 → 00 → 01 → 00 ... $50000/100000 = 50\%$

01 → 00 → 01 → 00 → 01 ... $50000/100000 = 50\%$

10 → 01 → 10 → 01 → 10 ... $100000/100000 = 100\%$

11 → 10 → 11 → 10 → 11 ... $50000/100000 = 50\%$

Advanced Predictors

- Leverage Correlation

```
if (x[i] < 7) then
    y += 1;
if (x[i] < 5) then
    c -= 4;
```

If first condition false,
second condition also false

```
if (aa == 2) then
    aa = 0;
if (bb == 2) then
    bb = 0;
if (aa != bb) {  
}
```

Outcome of third branch depends on
previous two branches

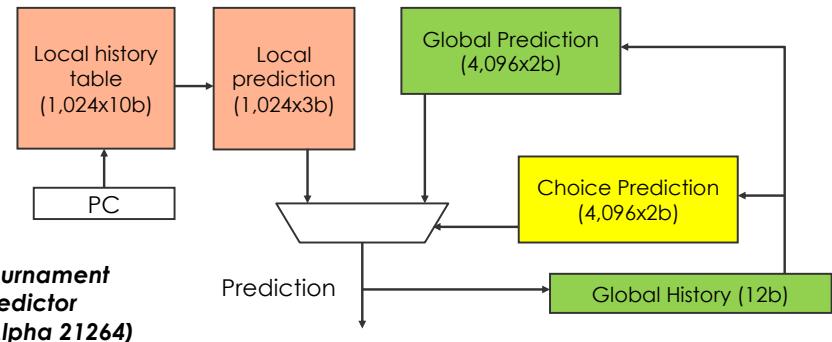
- Learn Optimal Prediction

- Perceptron-based Neural Network

Branch Predictors in Modern Processors

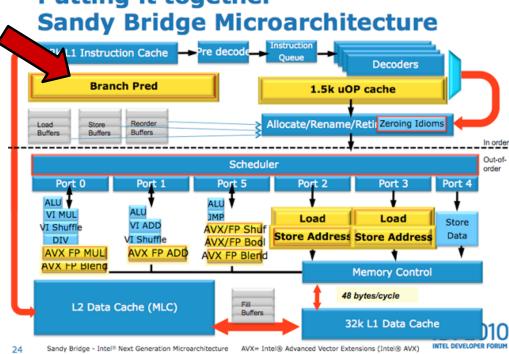
- No predictor is clearly the best

- Design a predictor to predict which predictor will predict better! ☺



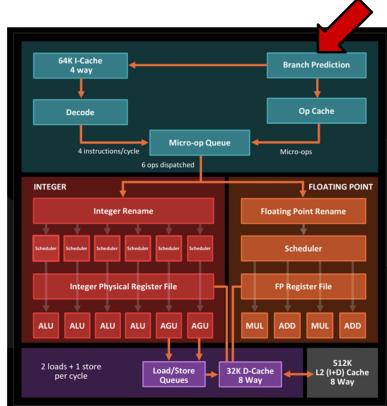
Branch Predictors in Modern Processors

Putting it together Sandy Bridge Microarchitecture



24 Sandy Bridge - Intel® Next Generation Microarchitecture AVX™ - Intel® Advanced Vector Extensions [Intel® AVX] INTEL DEVELOPER forum

Intel Sandy Bridge



AMD Ryzen